

Importing Library

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

Importing Dataset

```
In [169]: df=pd.read_csv(r"C:\Users\Shubham\Desktop\Data Science\Data Science Class\Stats a
df.head()
```

```
Out[169]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase
0	65	Female	0.7	0.1	187	16
1	62	Male	10.9	5.5	699	64
2	62	Male	7.3	4.1	490	60
3	58	Male	1.0	0.4	182	14
4	72	Male	3.9	2.0	195	27

Data Discription

```
In [170]: df.Dataset.value_counts()
```

```
Out[170]: 1    416
2     167
Name: Dataset, dtype: int64
```

In [171]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 583 entries, 0 to 582
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   583 non-null    int64
1   Gender                               583 non-null    object
2   Total_Bilirubin                      583 non-null    float64
3   Direct_Bilirubin                    583 non-null    float64
4   Alkaline_Phosphotase                583 non-null    int64
5   Alamine_Aminotransferase            583 non-null    int64
6   Aspartate_Aminotransferase          583 non-null    int64
7   Total_Protiens                      583 non-null    float64
8   Albumin                             583 non-null    float64
9   Albumin_and_Globulin_Ratio          579 non-null    float64
10  Dataset                             583 non-null    int64
dtypes: float64(5), int64(5), object(1)
memory usage: 50.2+ KB
```


Checking Null Values

In [172]: df.isna().sum()

```
Out[172]: Age                                0
Gender                                0
Total_Bilirubin                        0
Direct_Bilirubin                      0
Alkaline_Phosphotase                  0
Alamine_Aminotransferase              0
Aspartate_Aminotransferase            0
Total_Protiens                        0
Albumin                              0
Albumin_and_Globulin_Ratio            4
Dataset                              0
dtype: int64
```

```
In [173]: def boxplots(col):
            sns.boxplot(df[col])
            plt.show()

            for i in list(df.select_dtypes(exclude=['object']).columns)[1:]:
                boxplots(i)
```



```
In [174]: df.Albumin_and_Globulin_Ratio=df.Albumin_and_Globulin_Ratio.fillna(df.Albumin_and_Globulin_Ratio.mean())
```

```
In [175]: df.isna().sum()
```

```
Out[175]: Age                0
Gender                0
Total_Bilirubin       0
Direct_Bilirubin      0
Alkaline_Phosphotase  0
Alamine_Aminotransferase  0
Aspartate_Aminotransferase  0
Total_Protiens        0
Albumin              0
Albumin_and_Globulin_Ratio  0
Dataset              0
dtype: int64
```

Encoding Concept

```
In [176]: df.Gender=df.Gender.astype('category')
            df.Gender=df.Gender.cat.codes
```

In [177]: `df.head()`

Out[177]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase
0	65	0	0.7	0.1	187	16
1	62	1	10.9	5.5	699	64
2	62	1	7.3	4.1	490	60
3	58	1	1.0	0.4	182	14
4	72	1	3.9	2.0	195	27

Splitting the data into independent variable and dependent variable

In []:

In [178]: `x=df.iloc[:,0:-1]`
`y=df.iloc[:, -1]`

In [179]: `x.head()`

Out[179]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase
0	65	0	0.7	0.1	187	16
1	62	1	10.9	5.5	699	64
2	62	1	7.3	4.1	490	60
3	58	1	1.0	0.4	182	14
4	72	1	3.9	2.0	195	27

In [180]: `y.head()`

Out[180]:

```
0    1
1    1
2    1
3    1
4    1
Name: Dataset, dtype: int64
```

Handling Imbalacne Data

In [181]: `import imblearn`

```
In [182]: from imblearn.over_sampling import RandomOverSampler
ros=RandomOverSampler()
x_sam,y_sam=ros.fit_resample(x,y)
```

```
In [183]: y_sam.value_counts()
```

```
Out[183]: 1    416
          2    416
          Name: Dataset, dtype: int64
```


```
In [184]: y.value_counts()
```

```
Out[184]: 1    416
          2    167
          Name: Dataset, dtype: int64
```

Data Preprocessing

```
In [ ]:
```

```
In [185]: plt.figure(figsize=(15,15))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.show()
```



Split the data into training and test for building the model and for prediction

In []:

```
In [186]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x_sam,y_sam,train_size=0.35,random
```

In []:

1. KNeighborsClassifier Model

```
In [187]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [188]: error_rate = []


for i in range(1,40):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train, y_train)
    pred_i = knn.predict(x_test)
    error_rate.append(np.mean(pred_i != y_test))
```

In [189]: error_rate

Out[189]: [0.2975970425138632,
0.34750462107208874,
0.32162661737523107,
0.35304990757855825,
0.35489833641404805,
0.36044362292051757,
0.3364140480591497,
0.3456561922365989,
0.3438077634011091,
0.3512014787430684,
0.32717190388170053,
0.3419593345656192,
0.34750462107208874,
0.3456561922365989,
0.32532347504621073,
0.31608133086876156,
0.32162661737523107,
0.3290203327171904,
0.32532347504621073,
0.34011090573012936,
0.3438077634011091,
0.3438077634011091,
0.3419593345656192,
0.3419593345656192,
0.34750462107208874,
0.3438077634011091,
0.35304990757855825,
0.34750462107208874,
0.34935304990757854,
0.3438077634011091,
0.3456561922365989,
0.3364140480591497,
0.3345656192236599,
0.32717190388170053,
0.3364140480591497,
0.34011090573012936,
0.3345656192236599,
0.3234750462107209,
0.3364140480591497]

Plotting Error Rate vs K_hyperparam_value Graph

```
In [190]: plt.figure(figsize=(16,16))
plt.plot(range(1,40), error_rate, color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title("Error Rate vs K_hyperparam_value")
plt.xlabel("K_HyperParam_Value")
plt.ylabel("Error Rate")
plt.show()
```



Finding the N valuse based on the above data

```
In [191]: knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(x_train, y_train)
```

```
Out[191]: KNeighborsClassifier(n_neighbors=1)
```


Predict the data

In []:

```
In [192]: y_pred_train = knn.predict(x_train)
y_pred_test = knn.predict(x_test)
```

Evaluate the model

In []:

```
In [193]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [194]: print(confusion_matrix(y_train, y_pred_train))
print("*****"*5)
print(confusion_matrix(y_test, y_pred_test))
```

```
[[146  0]
 [ 0 145]]
*****
*
[[178  92]
 [ 69 202]]
```

```
In [195]: print(classification_report(y_train, y_pred_train))
print("*****"*5)
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	146
2	1.00	1.00	1.00	145
accuracy			1.00	291
macro avg	1.00	1.00	1.00	291
weighted avg	1.00	1.00	1.00	291

```
*****
*
```

	precision	recall	f1-score	support
1	0.72	0.66	0.69	270
2	0.69	0.75	0.72	271
accuracy			0.70	541
macro avg	0.70	0.70	0.70	541
weighted avg	0.70	0.70	0.70	541

```
In [196]: print("Training Accuracy :", accuracy_score(y_train, y_pred_train))
print("*****"*5)
print("Test Accuracy :", accuracy_score(y_test, y_pred_test))
```

Training Accuracy : 1.0

*

Test Accuracy : 0.7024029574861368

Cross Validation approach - K-Fold Method

In []:

```
In [197]: from sklearn.model_selection import cross_val_score
training_accuracy = cross_val_score(knn, x_train, y_train, cv=10)
test_accuracy = cross_val_score(knn, x_test, y_test, cv=10)
print("Training Accuracy after CV :", training_accuracy.mean())
print("*****"*5)
print("Test Accuracy after CV :", test_accuracy.mean())
```

Training Accuracy after CV : 0.7222988505747128

Test Accuracy after CV : 0.7984175084175085

In []:

2. BaggingClassifier Model

```
In [198]: from sklearn.ensemble import BaggingClassifier
bagging=BaggingClassifier()
bagging.fit(x_train,y_train)
```

Out[198]: BaggingClassifier()

Predict the data

In []:

```
In [199]: y_pred_train_bgg=bagging.predict(x_train)
y_pred_test_bgg=bagging.predict(x_test)
```

Evaluate the model

In []:

In [200]: `from sklearn.metrics import confusion_matrix, classification_report, accuracy_score`In [201]: `print(confusion_matrix(y_train,y_pred_train_bgg))
print(confusion_matrix(y_test,y_pred_test_bgg))`

```
[[145  1]
 [ 1 144]]
[[194  76]
 [ 45 226]]
```

In [202]: `print(classification_report(y_train,y_pred_train_bgg))
print(classification_report(y_test,y_pred_test_bgg))`

	precision	recall	f1-score	support
1	0.99	0.99	0.99	146
2	0.99	0.99	0.99	145
accuracy			0.99	291
macro avg	0.99	0.99	0.99	291
weighted avg	0.99	0.99	0.99	291

	precision	recall	f1-score	support
1	0.81	0.72	0.76	270
2	0.75	0.83	0.79	271
accuracy			0.78	541
macro avg	0.78	0.78	0.78	541
weighted avg	0.78	0.78	0.78	541

In [203]: `print(accuracy_score(y_train,y_pred_train_bgg))
print(accuracy_score(y_test,y_pred_test_bgg))`

```
0.993127147766323
0.7763401109057301
```

Cross Validation approach - K-Fold Method

In []:

```
In [204]: training_accuracy = cross_val_score(bagging, x_train, y_train, cv=10)
test_accuracy = cross_val_score(bagging, x_test, y_test, cv=10)
print("Training Accuracy after CV :", training_accuracy.mean())
print("*****5)
print("Test Accuracy after CV :", test_accuracy.mean())
```

Training Accuracy after CV : 0.7185057471264369

Test Accuracy after CV : 0.8188888888888888

In []:

3. RandomForestClassifier Model

```
In [205]: from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=200,criterion='entropy',bootstrap=True,oo
rf.fit(x_train,y_train)
```

Out[205]: RandomForestClassifier(criterion='entropy', n_estimators=200)

Predict the data

```
In [206]: y_pred_train_rf=rf.predict(x_train)
y_pred_test_rf=rf.predict(x_test)
```

Evaluate the model

In []:

```
In [207]: print(confusion_matrix(y_train,y_pred_train_rf))
print(confusion_matrix(y_test,y_pred_test_rf))
```

```
[[146  0]
 [ 0 145]]
[[180  90]
 [ 48 223]]
```

```
In [208]: print(classification_report(y_train,y_pred_train_rf))
print(classification_report(y_test,y_pred_test_rf))
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	146
2	1.00	1.00	1.00	145
accuracy			1.00	291
macro avg	1.00	1.00	1.00	291
weighted avg	1.00	1.00	1.00	291

	precision	recall	f1-score	support
1	0.79	0.67	0.72	270
2	0.71	0.82	0.76	271
accuracy			0.74	541
macro avg	0.75	0.74	0.74	541
weighted avg	0.75	0.74	0.74	541

```
In [209]: print(accuracy_score(y_train,y_pred_train_rf))
print(accuracy_score(y_test,y_pred_test_rf))
```

```
1.0
0.744916820702403
```

Cross Validation approach - K-Fold Method

```
In [ ]:
```

```
In [210]: training_accuracy = cross_val_score(rf, x_train, y_train, cv=10)
test_accuracy = cross_val_score(rf, x_test, y_test, cv=10)
print("Training Accuracy after CV :", training_accuracy.mean())
print("*****"*5)
print("Test Accuracy after CV :", test_accuracy.mean())
```

```
Training Accuracy after CV : 0.7357471264367816
*****
*****
Test Accuracy after CV : 0.8040404040404041
```

```
In [ ]:
```

```
In [ ]:
```

